

w2_assessment

December 4, 2021

In this notebook, we'll ask you to find numerical summaries for a certain set of data. You will use the values of what you find in this assignment to answer questions in the quiz that follows (we've noted where specific values will be requested in the quiz, so that you can record them.)

We'll also ask you to create some of the plots you have seen in previous lectures.

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import scipy.stats as stats
%matplotlib inline
import matplotlib.pyplot as plt
pd.set_option('display.max_columns', 100)
```

```
path = "nhanes_2015_2016.csv"
```

```
In [2]: # First, you must import the data from the path given above
df = pd.read_csv(path)
```

```
In [6]: # Next, look at the 'head' of our DataFrame 'df'.
df.head(10)
```

*# If you can't remember a function, open a previous notebook or video as a reference
or use your favorite search engine to look for a solution*

```
Out[6]:
```

	SEQN	ALQ101	ALQ110	ALQ130	SMQ020	RIAGENDR	RIDAGEYR	RIDRETH1	\
0	83732	1.0	NaN	1.0	1	1	62	3	
1	83733	1.0	NaN	6.0	1	1	53	3	
2	83734	1.0	NaN	NaN	1	1	78	3	
3	83735	2.0	1.0	1.0	2	2	56	3	
4	83736	2.0	1.0	1.0	2	2	42	4	
5	83737	2.0	2.0	NaN	2	2	72	1	
6	83741	1.0	NaN	8.0	1	1	22	4	
7	83742	1.0	NaN	1.0	2	2	32	1	
8	83743	NaN	NaN	NaN	2	1	18	5	
9	83744	1.0	NaN	NaN	2	1	56	4	

	DMDCITZN	DMDDEDUC2	DMDMARTL	DMDHHSIZ	WTINT2YR	SDMVPSU	SDMVSTRA	\
0	1.0	5.0	1.0	2	134671.37	1	125	

1	2.0	3.0	3.0	1	24328.56	1	125
2	1.0	3.0	1.0	2	12400.01	1	131
3	1.0	5.0	6.0	1	102718.00	1	131
4	1.0	4.0	3.0	5	17627.67	2	126
5	2.0	2.0	4.0	5	11252.31	1	128
6	1.0	4.0	5.0	3	37043.09	2	128
7	2.0	4.0	1.0	4	22744.36	1	125
8	1.0	NaN	NaN	3	18526.16	2	122
9	1.0	3.0	3.0	1	20395.54	2	126

	INDFMPIR	BPXSY1	BPXDI1	BPXSY2	BPXDI2	BMXWT	BMXHT	BMXBMI	BMXLEG	\
0	4.39	128.0	70.0	124.0	64.0	94.8	184.5	27.8	43.3	
1	1.32	146.0	88.0	140.0	88.0	90.4	171.4	30.8	38.0	
2	1.51	138.0	46.0	132.0	44.0	83.4	170.1	28.8	35.6	
3	5.00	132.0	72.0	134.0	68.0	109.8	160.9	42.4	38.5	
4	1.23	100.0	70.0	114.0	54.0	55.2	164.9	20.3	37.4	
5	2.82	116.0	58.0	122.0	58.0	64.4	150.0	28.6	34.4	
6	2.08	110.0	70.0	112.0	74.0	76.6	165.4	28.0	38.8	
7	1.03	120.0	70.0	114.0	70.0	64.5	151.3	28.2	34.1	
8	5.00	NaN	NaN	NaN	NaN	72.4	166.1	26.2	NaN	
9	1.19	178.0	116.0	180.0	114.0	108.3	179.4	33.6	46.0	

	BMXARML	BMXARMC	BMXWAIST	HIQ210
0	43.6	35.9	101.1	2.0
1	40.0	33.2	107.9	NaN
2	37.0	31.0	116.5	2.0
3	37.7	38.3	110.1	2.0
4	36.0	27.2	80.4	2.0
5	33.5	31.4	92.9	NaN
6	38.0	34.0	86.6	NaN
7	33.1	31.5	93.3	2.0
8	NaN	NaN	NaN	2.0
9	44.1	38.5	116.0	2.0

How many rows can you see when you don't put an argument into the previous method? 5
 How many rows can you see if you use an int (n) as an argument? n rows Can you use a float as an argument? No.

```
In [7]: # Lets only consider the feature (or variable) 'BPXSY2' (Systolic Blood pressure, meas
bp = df['BPXSY2']
```

0.1 Numerical Summaries

0.1.1 Find the mean (note this for the quiz that follows)

```
In [10]: # What is the mean of 'BPXSY2'?
bp_mean = bp.mean()
bp_mean
```

```
Out[10]: 124.78301716350497
```

In the method you used above, how are the rows of missing data treated? Are the excluded entirely? Are they counted as zeros? Something else? By default, they are excluded (NA/null values) when computing the result.

If you used a library function, try looking up the documentation using the code:

```
help(function_you_used)
```

For example:

```
help(np.sum)
```

```
In [16]: help(df.mean)
```

Help on method mean in module pandas.core.frame:

```
mean(axis=None, skipna=None, level=None, numeric_only=None, **kwargs) method of pandas.core.frame
```

Return the mean of the values for the requested axis.

Parameters

axis : {index (0), columns (1)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

numeric_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

**kwargs

Additional keyword arguments to be passed to the function.

Returns

mean : Series or DataFrame (if level specified)

.dropna() To make sure we know that we aren't treating missing data in ways we don't want, let's go ahead and drop all the nans from our Series 'bp'

```
In [17]: bp = bp.dropna()
```

```
In [23]: bp.head()
```

```
Out [23]: 0    124.0
          1    140.0
          2    132.0
          3    134.0
          4    114.0
          Name: BPXSY2, dtype: float64
```

0.1.2 Find the:

- Median
- Max
- Min
- Standard deviation
- Variance

You can implement any of these from base python (that is, without any of the imported packages), but there are simple and intuitively named functions in the numpy library for all of these. You could also use the fact that 'bp' is not just a list, but is a pandas.Series. You can find pandas.Series attributes and methods [here](#)

A large part of programming is being able to find the functions you need and to understand the documentation formatting so that you can implement the code yourself, so we highly encourage you to search the internet whenever you are unsure!

```
In [20]: bp.describe()
```

```
Out [20]: count      5535.000000
          mean       124.783017
          std        18.527012
          min        84.000000
          25%        112.000000
          50%        122.000000
          75%        134.000000
          max        238.000000
          Name: BPXSY2, dtype: float64
```

```
In [21]: variance = bp.std()**2
          variance
```

```
Out [21]: 343.25016328394815
```

```
In [45]: # quartis
          quartil_25 = bp.quantile(q=0.25)
          quartil_50 = bp.quantile(q=0.5)
          quartil_75 = bp.quantile(q=0.75)
          print(quartil_25, quartil_50, quartil_75)
```

```
112.0 122.0 134.0
```

```
In [46]: # intervalor interquartil
        iqr = quartil_75 - quartil_25
        iqr
```

```
Out[46]: 22.0
```

0.1.3 Example:

Find the difference of an element in 'bp' compared with the previous element in 'bp'.

```
In [22]: # Using the fact that 'bp' is a pd.Series object, can use the pd.Series method diff()
        # call this method by: pd.Series.diff()
        diff_by_series_method = bp.diff()
        # note that this returns a pd.Series object, that is, it had an index associated with
        diff_by_series_method.values # only want to see the values, not the index and values
```

```
Out[22]: array([ nan,  16.,  -8., ...,  30., -40.,   8.])
```

```
In [24]: # Now use the numpy library instead to find the same values
        # np.diff(array)
        diff_by_np_method = np.diff(bp)
        diff_by_np_method
        # note that this returns an 'numpy.ndarray', which has no index associated with it, a
        # the nan we get by the Series method
```

```
Out[24]: array([ 16.,  -8.,   2., ...,  30., -40.,   8.])
```

```
In [29]: # We could also implement this ourselves with some looping
        diff_by_me = [] # create an empty list
        for i in range(len(bp.values)-1): # iterate through the index values of bp
            diff = bp.values[i+1] - bp.values[i] # find the difference between an element and
            diff_by_me.append(diff) # append to out list
        np.array(diff_by_me) # format as an np.array
```

```
Out[29]: array([ 16.,  -8.,   2., ...,  30., -40.,   8.])
```

0.1.4 Your turn (note these values for the quiz that follows)

```
In [30]: bp_median = bp.median()
        bp_median
```

```
Out[30]: 122.0
```

```
In [31]: bp_max = bp.max()
        bp_max
```

```
Out[31]: 238.0
```

```
In [32]: bp_min = bp.min()
        bp_min
```

```
Out [32]: 84.0
```

```
In [33]: bp_std = bp.std()  
bp_std
```

```
Out [33]: 18.527011720294997
```

```
In [34]: bp_var = bp.var()  
bp_var
```

```
Out [34]: 343.2501632839482
```

0.1.5 How to find the interquartile range (note this value for the quiz that follows)

This time we need to use the `scipy.stats` library that we imported above under the name 'stats'

```
In [44]: bp_iqr = stats.iqr(bp)  
bp_iqr
```

```
Out [44]: 22.0
```


0.2 Visualizing the data

Next we'll use what you have learned from the *Tables, Histograms, Boxplots in Python* video

```
In [48]: # use the Series.describe() method to see some descriptive statistics of our Series 'bp'  
bp_descriptive_stats = bp.describe()  
bp_descriptive_stats
```

```
Out [48]: count      5535.000000  
mean       124.783017  
std        18.527012  
min         84.000000  
25%        112.000000  
50%        122.000000  
75%        134.000000  
max        238.000000  
Name: BPXSY2, dtype: float64
```

```
In [50]: # Make a histogram of our 'bp' data using the seaborn library we imported as 'sns'  
sns.distplot(bp, kde=False).set_title("BP Histogram")  
plt.show()
```



Is your histogram labeled and does it have a title? If not, try appending

```
.set(title='your_title', xlabel='your_x_label', ylabel='your_y_label')
```

or just

```
.set(title='your_title')
```

to your graphing function

```
In [52]: # Make a boxplot of our 'bp' data using the seaborn library. Make sure it has a title
sns.distplot(bp, kde=False).set(title="BP Histogram", ylabel='Frequency', xlabel="BPXSY2")
plt.show()
```

